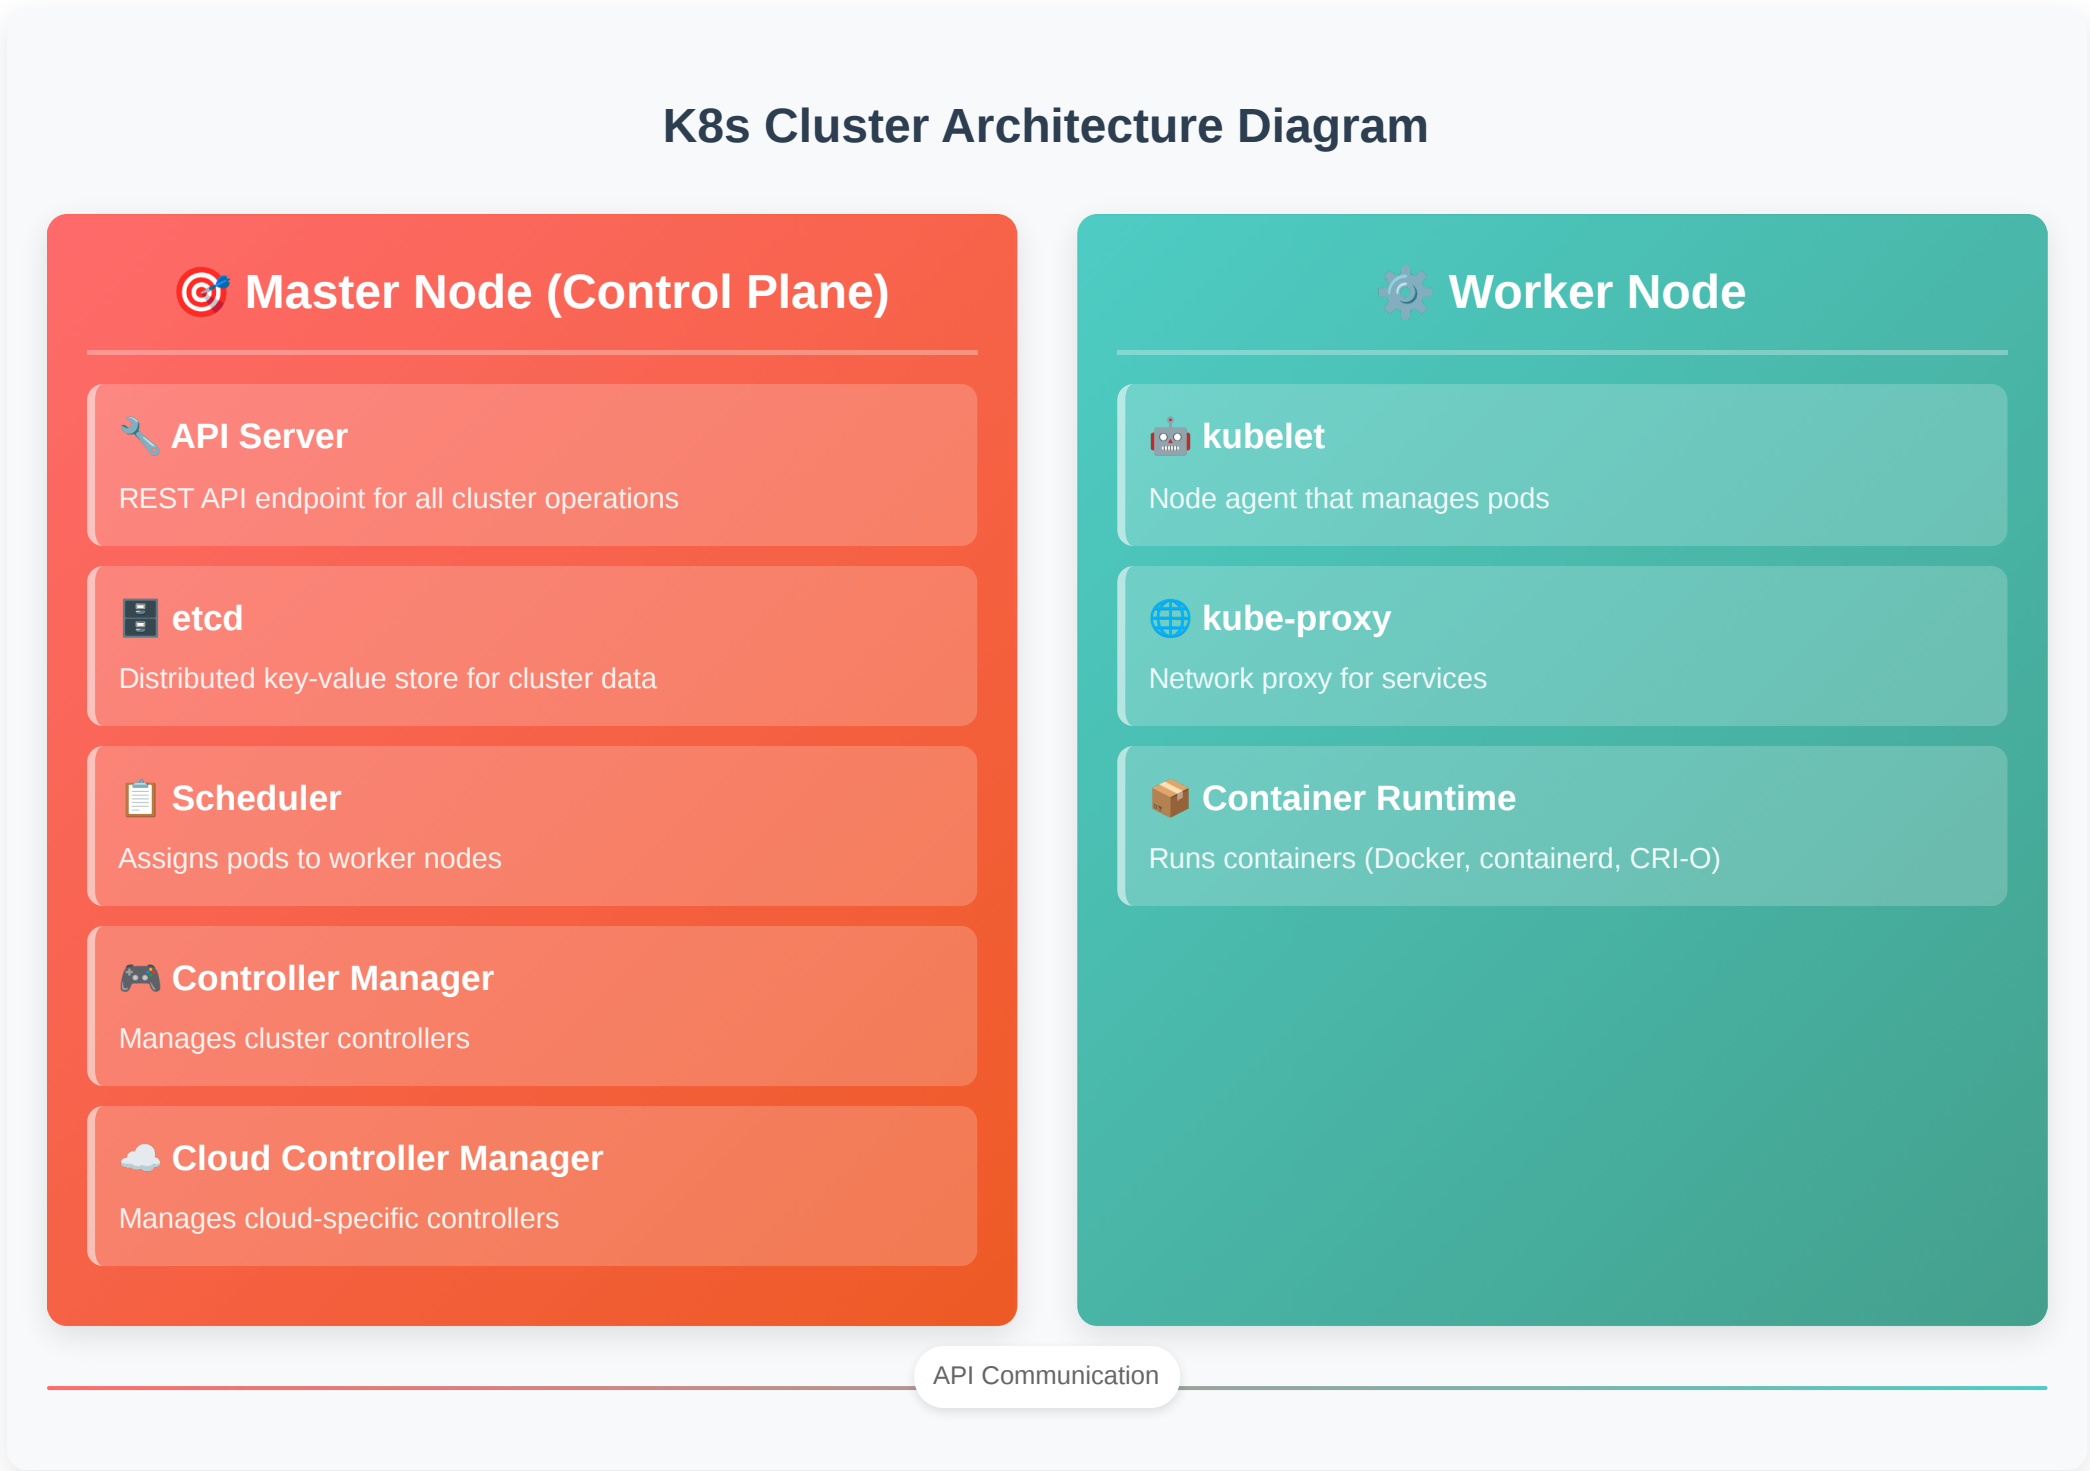




Kubernetes Cluster Architecture



Master Node Components (Control Plane)

API Server (kube-apiserver)

The **API Server** is the central management entity and the only component that directly interacts with etcd. It:

- Exposes the Kubernetes REST API
- Validates and processes API requests
- Authenticates and authorizes requests
- Acts as the frontend for the cluster's shared state

etcd

A highly available **distributed key-value store** that:

- Stores all cluster data and configuration
- Maintains the desired state of the cluster
- Provides strong consistency and reliability
- Supports backup and restore operations

Scheduler (kube-scheduler)

The **Scheduler** is responsible for:

- Selecting optimal nodes for pod placement
- Considering resource requirements and constraints
- Evaluating node affinity and anti-affinity rules
- Balancing workloads across the cluster

Controller Manager (kube-controller-manager)

Runs multiple **controllers** including:

- Deployment Controller:** Manages ReplicaSets for deployments
- ReplicaSet Controller:** Ensures desired number of pod replicas
- Node Controller:** Monitors node health and status
- Service Controller:** Manages service endpoints

Cloud Controller Manager (CCM)

Manages **cloud-specific controllers** for:

- Load balancers and external IPs
- Storage volumes and persistent volumes
- Node lifecycle management in cloud environments
- Integration with cloud provider APIs

Worker Node Components

kubelet

The **node agent** that:

- Communicates with the API server
- Manages pods and containers on the node
- Reports node and pod status back to the control plane
- Performs health checks on containers
- Mounts volumes and handles pod networking

kube-proxy

A **network proxy** that:

- Implements Kubernetes service abstraction
- Manages network rules for service discovery
- Handles load balancing for services
- Maintains network connectivity between pods
- Supports different proxy modes (iptables, IPVS, userspace)

Container Runtime

The **container runtime** options include:

- Docker:** Most popular container runtime
- containerd:** Lightweight, high-performance runtime
- CRI-O:** OCI-compliant runtime specifically for Kubernetes
- Responsibilities:** Pulling images, running containers, managing container lifecycle

How Components Work Together

Pod Creation Flow:

- User** submits a pod manifest to the **API Server**
- API Server** validates and stores the pod spec in **etcd**
- Scheduler** watches for unscheduled pods and assigns them to nodes
- kubelet** on the selected node pulls the pod spec and creates containers
- Container Runtime** pulls images and starts containers
- kube-proxy** updates network rules for the new pod

Service Discovery & Load Balancing:

- Services** are created via API Server and stored in etcd
- kube-proxy** watches for service changes and updates iptables rules
- DNS** resolution is handled by CoreDNS (typically)
- Load balancing** is performed by kube-proxy across healthy endpoints

Key Takeaways

- Master Node:** Manages the cluster state and makes scheduling decisions
- Worker Nodes:** Run the actual application workloads
- API Server:** Central hub for all cluster communication
- etcd:** Single source of truth for cluster state
- kubelet:** Ensures pods are running as expected on each node
- kube-proxy:** Handles network connectivity and service discovery